

2 - Image filtering

Spatial filtering

Pixel operator: Computes an output value at each pixel location, based on the input pixel value

- Transform pixel based on its value
- Gamma correction

Local operator: Computes an output value at each pixel location, based on a neighbourhood of pixels around the input pixel

- Transform pixel based on its neighbours
- e.g. sharpening filter

Linear filtering: Output pixel's value is a weighted sum of a neighbourhood around the input pixel

Cross-correlation vs. Convolution

$$g(i, j) = h(u, v) * f(i, j)$$

Output image g Kernel h Input image f

Convolution operator

$$g(i, j) = \sum_{u, v} f(i - u, j - v) h(u, v)$$

$$g(i, j) = h(u, v) \otimes f(i, j)$$

Output image g Kernel h Input image f

Cross-correlation convolution

$$g(i, j) = \sum_{u, v} f(i + u, j + v) h(u, v)$$

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

$F[x,y]$

a	b	c
d	e	f
g	h	i

$H[u,v]$

		a	b	c		
		d	e	f		
		g	h	i		

$F[x,y] * H[u,v]$

		i	h	g		
		f	e	d		
		c	b	a		

$F[x,y] \otimes H[u,v]$

- Cross-correlation: overlay filter on image
- Convolution: flip filter horizontally and vertically
- They are operations that apply a linear filter to an image.
- Illustration: L2.1 P13-18

Common filters

- Average/blur filters: average pixel values, blur the image
- Sharpening filters: subtract pixel from surround, increase fine detail
- Edge filters: compute difference between pixels, detect oriented edges in image

0	0	0
0	1	0
0	0	0

No change. Behave like a pixel operator

0	0	0
0	0	1
0	0	0

Shift left by 1 pixel

0	0	0
0	2	0
0	0	0

—


$\frac{1}{9}$

1	1	1
1	1	1
1	1	1


Sharpening filter. Accentuates differences

with local average.

Gaussian



$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$



Gaussian kernel
Kernel size: 5 x 5 px
 $\sigma = 1$

Sobel: L2.1 P26, 27

- Detect edges (Vertical, Horizontal)
- More weight on center which provides a bit smoothing.

Colour image filtering: Multiple channels

- Convolve each layer: 2D convolution in each colour channel -> Output is 3 channels
- 3D kernel: Output is 1 channel

Filter design examples

- Diagonal edges: e.g. $[[0, 1, 2], [-1, 0, 1], [-2, -1, 0]]$
- Simulate (linear) motion blur: e.g. 1D gaussian layer

Filters in practice

Properties of linear filters

- Commutative: $f * h = h * f$
 - Theoretically, no difference between kernel and image
 - But most implementations do care about order
- Associative: $(f * h1) * h2 = f * (h1 * h2)$
 - Usually one option is faster than the other – allows for more efficient implementations
- Distributive over addition
 - $f * (h1 + h2) = (f * h1) + (f * h2)$
- Multiplication cancels out
 - $kf * h = f * kh = k(f * h)$

Efficient filtering

- Multiple filters: generally more efficient to combine 2D filters ($h1 * h2 * h3 \dots$) and filter image just once L2.1 P35
- Separable filters: generally more efficient to filter with two 1D filters than one 2D filter L2.1 P37
- For example, the 2D Gaussian can be expressed as a product of two 1D Gaussians (in x and

$$G_{\sigma}(x, y) = \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2 + y^2}{2\sigma^2}}$$

$$= \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{x^2}{2\sigma^2}} \right) \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{y^2}{2\sigma^2}} \right)$$

Convolution output size

Valid convolution: The output image is smaller than the input image

Border handling L2.1 P39-42

- Pad with constant value
- Wrap image
- Clamp/replicate the border value
- Reflect image


Practical considerations

- Think about how to implement filters efficiently
 - Images are big, so efficient filtering can save a lot of time.
- Think about how to handle borders
 - No one-size-fits-all solution
 - Wrap is ideal for tiling textures (but not photos)
 - Clamp/replicate tends to work well for photos
- Linear filters: first step of almost all computer vision systems
- Linear filters are just a first step: you can't build complex feature detectors from just linear filters.

Frequency filtering

Fourier analysis (1D)

Any signals or pattern can be described as a sum of *sinusoids* (L2.2 P7-13).



Fourier transform

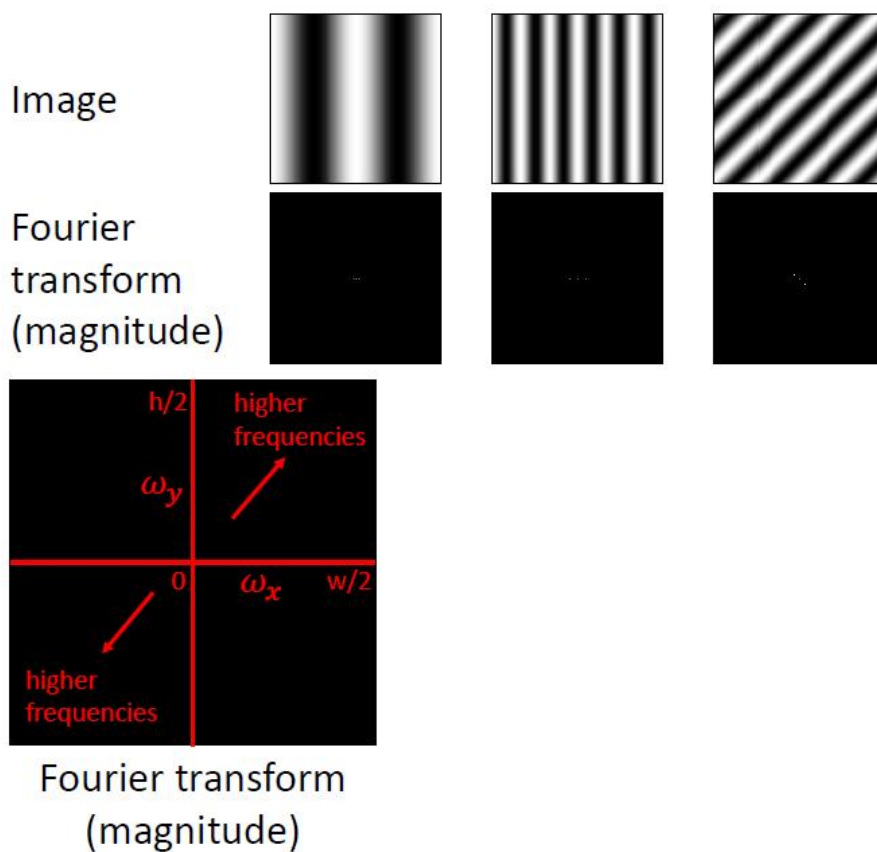
- Fourier transform decomposes signal into component frequencies
 - Values are complex numbers representing amplitude and phase of sinusoids
 - Time domain -> frequency domain (or, for images, spatial domain -> frequency domain)
- $F(\omega) = \int_{-\infty}^{\infty} f(x) e^{-2i\pi\omega x} dx$
- `scipy.fft` (1D), `scipy.fft2` (2D), `scipy.fftn` (3D+)
- Inverse Fourier transform converts from frequency domain back to space domain

Frequency Spectrum L2.2 P19-21

- Values in frequency domain are complex numbers
- For each frequency : magnitude (= amplitude) and angle (= phase)

Fourier analysis (images) L2.2 P24-34

- Any image can be represented by its Fourier transform
- Fourier transform = for each frequency, magnitude (amplitude) + phase



Magnitude and phase

- Magnitude is easy to read, giving the characteristics of the texture of the image.
- Phase represents actual structures and edges.
- Magnitude captures the holistic “texture” of an image, but the edges are mainly represented by Fourier phase

Frequency filtering

- Operations in the spatial domain have equivalent operations in frequency domain
- Convolution in spatial domain = multiplication in frequency domain

$$FT[h * f] = FT[h]FT[f]$$

- Inverse:

$$FT^{-1}[hf] = FT^{-1}[h] * FT^{-1}[f]$$

Band pass filter: A filter that removes a range of frequencies from a signal

- **Low pass filter**
 - Keep low spatial frequencies, remove high frequencies L2.2 P38-40
 - Equivalent to blurring the image
- **High pass filter**
 - Keep high spatial frequencies, remove low frequencies L2.2 P41-43
- "Ringing problem": Proof by inverse convolution theorem
 - Use Gaussian low/high pass filter

Summary

- Images can be filtered in the spatial domain, or the frequency domain
- Operations in one domain have an equivalent in the other domain
 - Convolution in spatial domain = multiplication in Fourier domain
- Modelling filters in both domains can help understand/debug what a filter is doing

Applications

Image compression

- Frequency domain is a convenient space for image compression
- Human visual system is not very sensitive to contrast in high spatial frequencies
- Discarding information in high spatial frequencies doesn't change the "look" of an image
- JPEG compression: break image into 8x8 pixel blocks, each represented in frequency space
- Discrete cosine transform (DCT)
- High spatial frequency components are quantised

Image forensic (鉴定) L2.2 P54-56

Summary

- Any image can be represented in either the spatial or the frequency domain
- Frequency domain is a convenient space for many applications:
 - Filtering
 - Compression
 - Forensics
 - Frequency-based features